

Non-Local Image Denoising & Integral Images

Slides by: Stephan Garbin





1. Non-Local Denoising

- Problem
- Noise Model
- Motivation for Non-Local Filters
- Non-Local Means
- Weighting for Non-Local Means

Outline



1. Non-Local Denoising

- Problem
- Noise Model
- Motivation for Non-Local Filters
- Non-Local Means
- Weighting for Non-Local Means

2. Efficient Template Matching

- Problem
- What is an Integral Image?
- How do we calculate it?
- How do we use it?



Non-Local Denoising

Our Task

Noisy Image

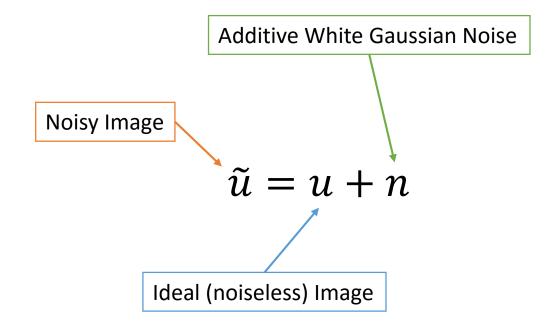
Denoised Image



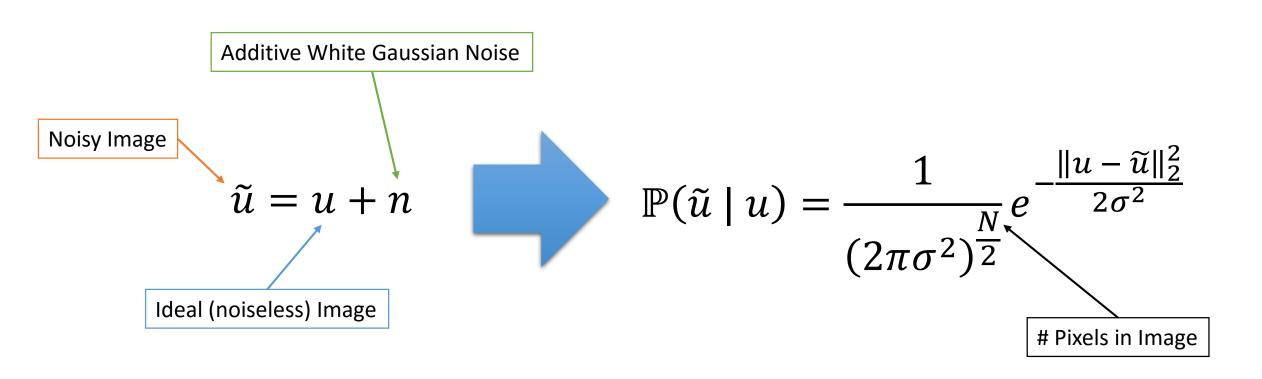


Noise Model

• Assumption: Additive, Zero-Mean, 'White' Gaussian Noise



Noise Model



Noise Model

Conditional Distribution

Noise Model

$$\tilde{u} = u + n$$

$$\mathbb{P}(\tilde{u} \mid u) = \frac{1}{(2\pi\sigma^2)^{\frac{N}{2}}} e^{\frac{\|u - \tilde{u}\|_2^2}{2\sigma^2}}$$

We want to recover u from \tilde{u}

How?

• (Gaussian) Low-Pass?



How?

- (Gaussian) Low-Pass?
- Bilateral Filter?



How?

- (Gaussian) Low-Pass?
- Bilateral Filter?

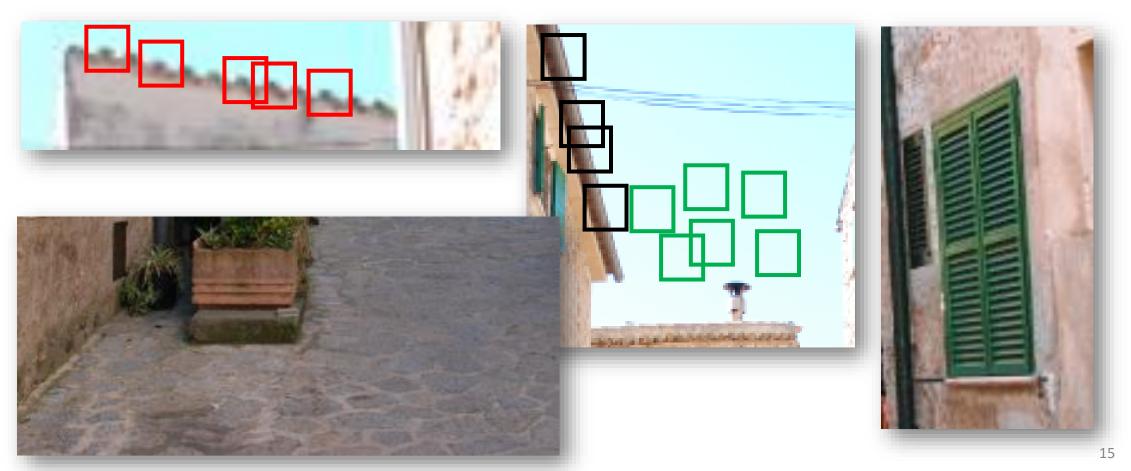


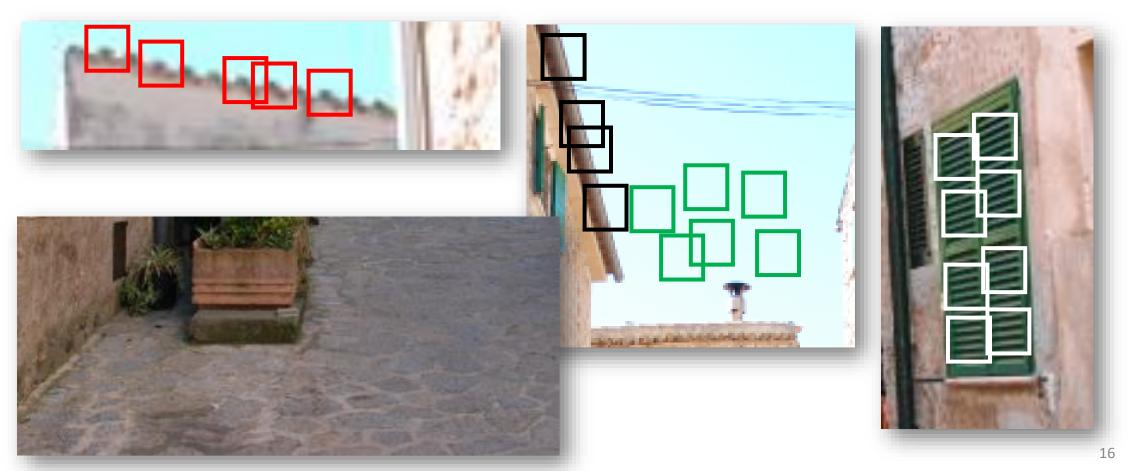
- Too much detail lost
- Blurry results

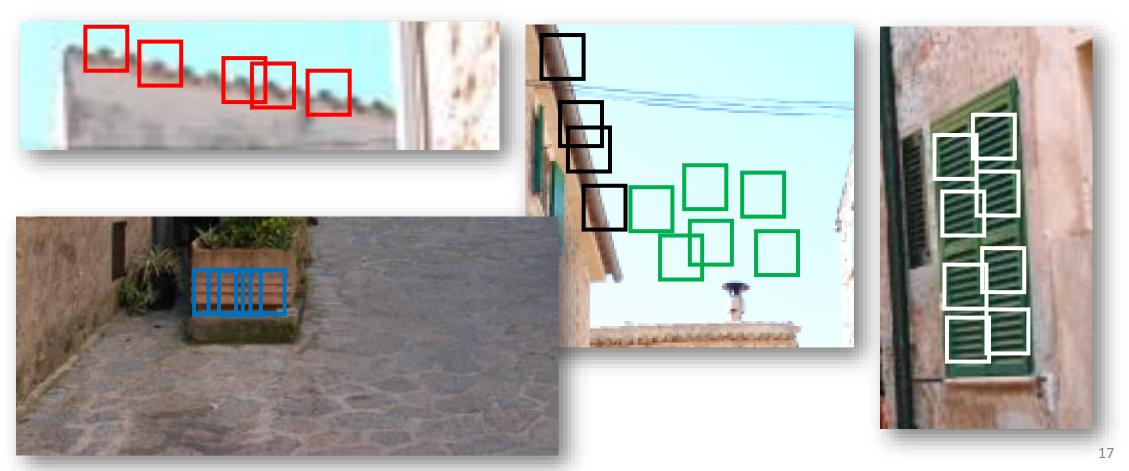












Non-Local Means: Concept

"Restore current pixel as a weighted average of all pixels in the image"

Non-Local Means: Concept

In Practise:

Local neighbourhood around pixel (e.g. 30x30)

"Restore current pixel as a weighted average of <u>all</u> pixels in the image"

Non-Local Means: Concept

In Practise:

Local neighbourhood around pixel (e.g. 30x30)

"Restore current pixel as a <u>weighted</u> average of <u>all</u> pixels in the image"

Similarity determined by *patches* around pixels

- Similar pixels count more
- Dissimilar pixels should not count

1. Block-Matching

- <u>Procedure:</u> For every reference patch, centred at every pixel, find k most similar patches
 - For NL-Means: All patches in local neighbourhood

1. Block-Matching

- <u>Procedure:</u> For every reference patch, centred at every pixel, find k most similar patches
 - For NL-Means: All patches in local neighbourhood
- *Result:* 3d block of 2d patches for each reference patch

Note:

Does not need to be stored explicitly for NL-Means

1. Block-Matching

- <u>Procedure:</u> For every reference patch, centred at every pixel, find k most similar patches
 - For NL-Means: All patches in local neighbourhood
- *Result:* 3d block of 2d patches for each reference patch

2. Filtering

Use 3d blocks to derive denoised estimate

1. Block-Matching

- <u>Procedure:</u> For every reference patch, centred at every pixel, find k most similar patches
 - For NL-Means: All patches in local neighbourhood
- Result: 3d block of 2d patches for each reference patch

2. Filtering

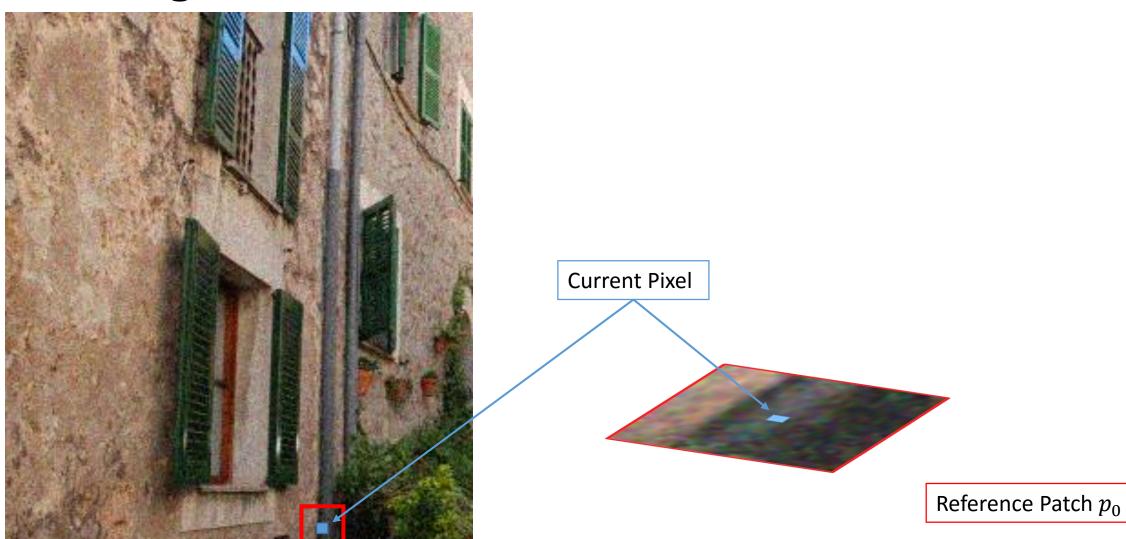
Use 3d blocks to derive denoised estimate

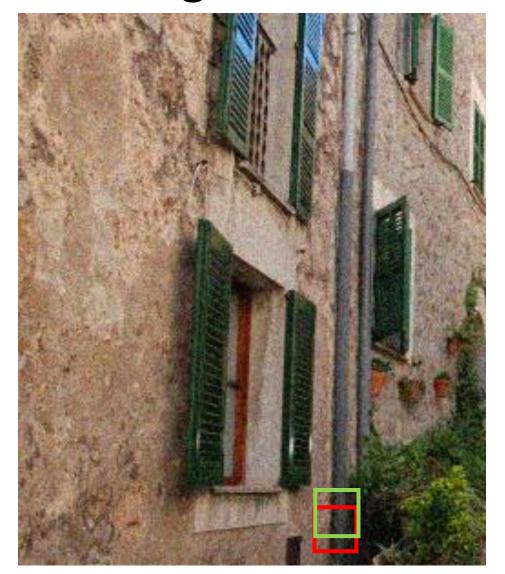
3. (Aggregation)

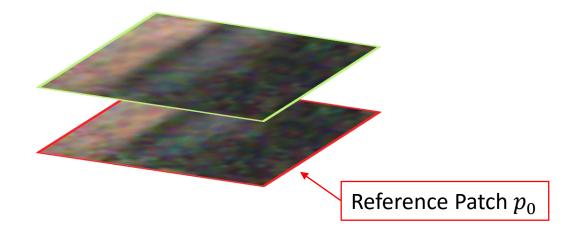
• Combine (overlapping) patches to form final image

Note:

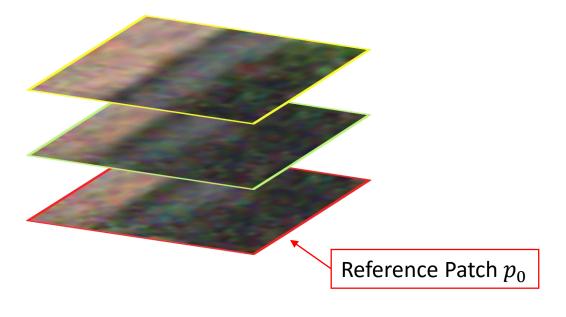
Only necessary for patch-wise implementation








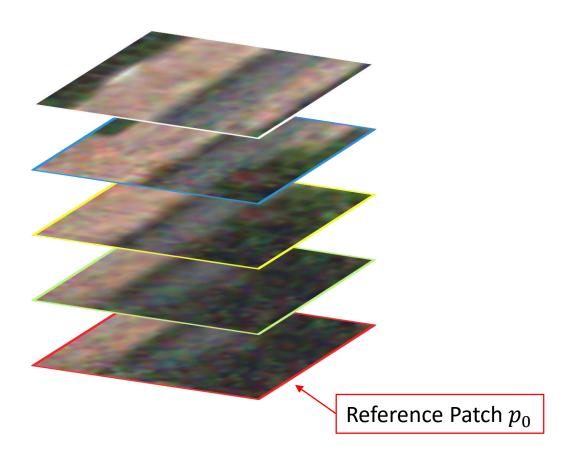







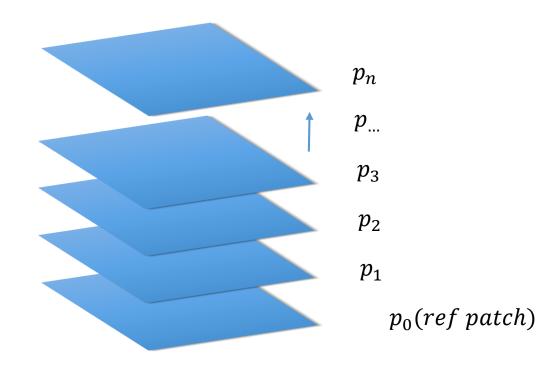









3D Block of Similar Patches




Pixel-Wise

Denoised pixel at location p
$$\hat{u}(p) = \frac{1}{C(p)} \sum_{q \in B(p,r)} \tilde{u}(q) \, w(p,q)$$

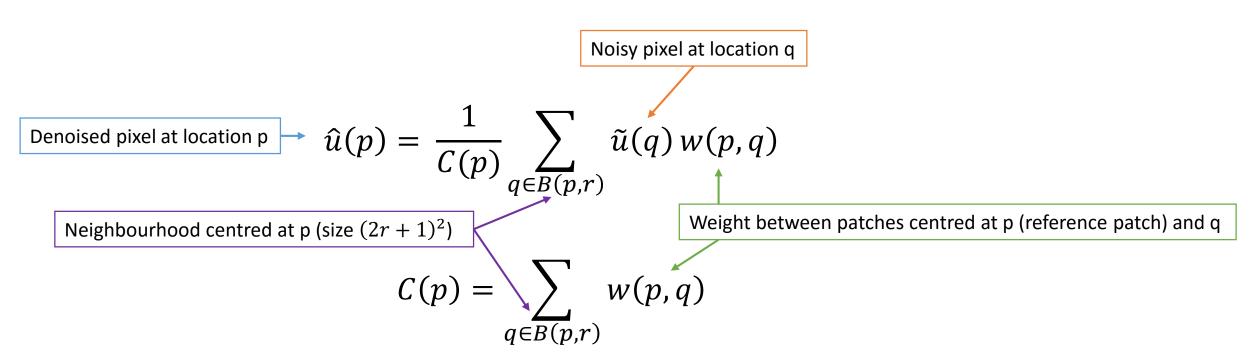
Pixel-Wise



Pixel-Wise



Pixel-Wise



Patch-Wise

- Use <u>all pixels of each patch</u> in the block for final estimate
- →Otherwise exactly the same as pixel-wise! ©

Patch-Wise

- Use <u>all pixels of each patch</u> in the block for final estimate
- →Otherwise exactly the same as pixel-wise! ©

Note:

- We now get multiple estimates per pixel → take average
- Reason: Patches centred at every pixel overlap

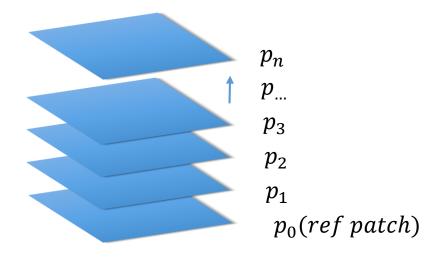
Filtering – Choosing the weights

Average pixels across all patches in a block?

Filtering – Choosing the weights

- Average pixels across all patches in a block?
- **Better:** Weight according to similarity to the reference patch

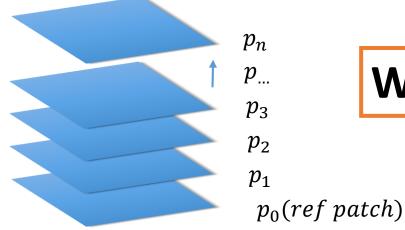
3D Block of Similar Patches



Filtering – Choosing the weights

- Average pixels across all patches in a block?
- **Better:** Weight according to similarity to the reference patch

3D Block of Similar Patches



We need a weight for each patch

- We want
 - Similar patches to count more
 - Allow more dissimilar patches to count in each block for strong noise
 - (Impose additional constraint on which patches to keep)

- We want
 - Similar patches to count more
 - Allow more dissimilar patches to count in each block for strong noise
 - (Impose additional constraint on which patches to keep)

- So far, we have assumed using all patches in a local neighbourhood
- However, we could exclude dissimilar patches from the start
 - \rightarrow We keep only a *subset* of patches in B(p,r)
- <u>Note:</u> Weighting function already does this to some degree (can be 0)

- We want
 - Similar patches to count more
 - Allow more dissimilar patches for strong noise
 - (Impose additional constraint on which patches to keep)

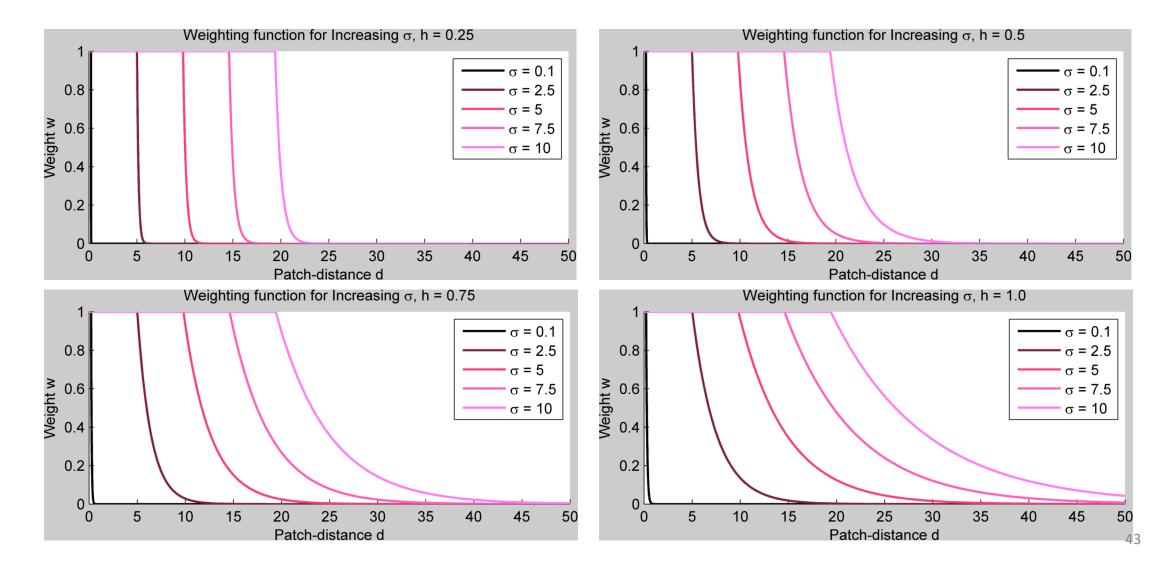
Squared Patch Distance (SSD)

Noise Standard Deviation

$$w(p,q) = e^{-\frac{\max(d^2-2\sigma^2,0)}{h^2}}$$

Weight between patches centred at p and q

Decay Parameter



Problems with Non-Local Methods

- Complexity (Can be quadratic or worse w.r.t. image size)
 - → Takes a long time to run
 - → In the real world, we deal with 2k, 4k, 6k Images

Problems with Non-Local Methods

- Complexity (Can be quadratic or worse w.r.t. image size)
 - → Takes a long time to run
 - → In the real world, we deal with 2k, 4k, 6k Images

Not all images are highly self-similar everywhere



Problems with Non-Local Methods

- Complexity (Can be quadratic or worse w.r.t. image size)
 - → Takes a long time to run
 - → In the real world, we deal with 2k, 4k, 6k Images

Not all images are highly self-similar everywhere



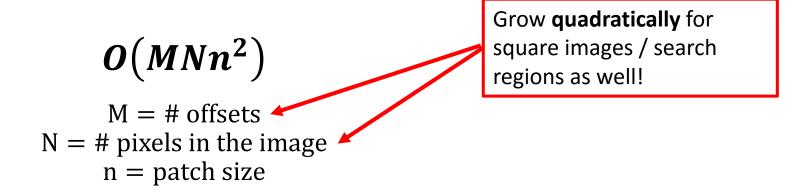
- Search **Database** of similar images
- Search other frames (if video)



Efficient Template Matching

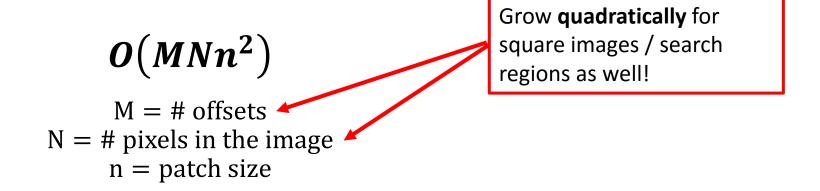
The Problem

 Calculating SSD for every patch centred at each pixel with every other patch in local neighbourhood is VERY expensive



The Problem

 Calculating SSD for every patch centred at each pixel with every other patch in local neighbourhood is VERY expensive





Integral Images can reduce number of arithmetic operations required to:

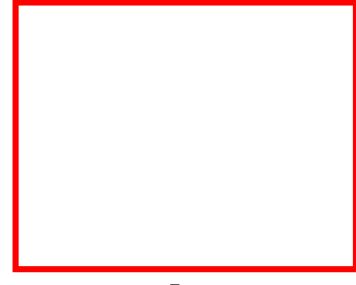
- The Integral Image I_{int} for an Image I stores at every pixel location,
 - The sum of all pixels in *I* to the left of the current location
 - And all pixels in *I* above of the current location
 - Including the pixel at the current location in *I*.

- The Integral Image I_{int} for an Image I stores at every pixel location,
 - The sum of all pixels in *I* to the left of the current location
 - And all pixels in *I* above of the current location
 - Including the pixel at the current location in *I*.

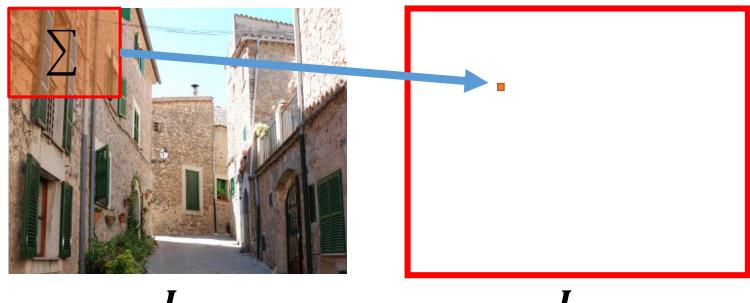
Integral Image location: x, y $I_{int}(x,y) = \sum_{p \leq x, q \leq y} I(p,q)$ Image location: p, q

- The Integral Image I_{int} for an Image I stores at every pixel location,
 - The sum of all pixels in *I* to the left of the current location
 - And all pixels in *I* above of the current location
 - Including the pixel at the current location in I.

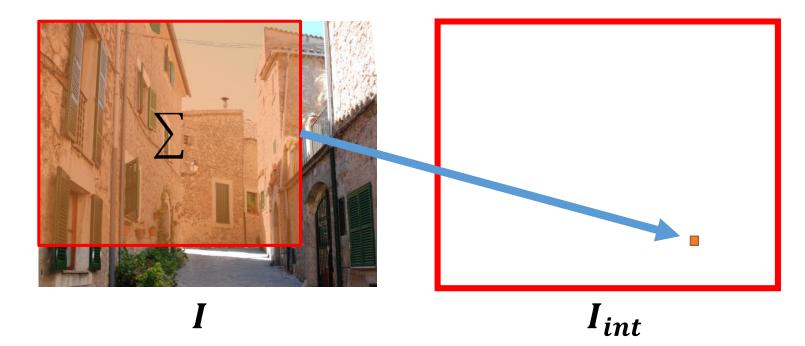




- The Integral Image I_{int} for an Image I stores at every pixel location,
 - The sum of all pixels in *I* to the left *of the current location*
 - And all pixels in *I* above of the current location
 - Including the pixel at the current location in *I*.

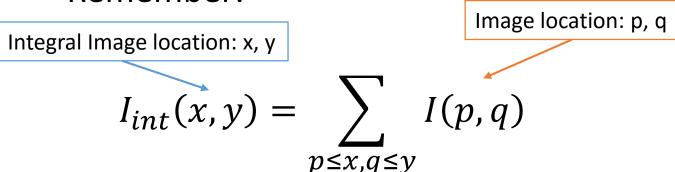


- The Integral Image I_{int} for an Image I stores at every pixel location,
 - The sum of all pixels in *I* to the left *of the current location*
 - And all pixels in *I* above of the current location
 - Including the pixel at the current location in *I*.



Computing it efficiently (serially)

• Remember:



Computing it efficiently (serially)

• Remember:

$$I_{int}(x,y) = \sum_{p \le x, q \le y} I(p,q)$$

• Use recurrence relation:

Cumulative sum (column-wise)

$$s(x,y) = s(x,y-1) + I(x,y)$$

$$I_{int}(x, y) = I_{int}(x - 1, y) + s(x, y)$$

Computing it efficiently (serially)

• Use recurrence relation:


Cumulative sum (column-wise)


$$s(x,y) = s(x,y-1) + I(x,y)$$


$$I_{int}(x,y) = I_{int}(x-1,y) + s(x,y)$$


Note Starting Point:

$$s(x,-1) = I_{int}(-1,y) = 0$$











 We can calculate the sum of all pixels in a patch in constant time, regardless of patch size!



Remember what pixels in I_{int} represent!



• We can build an integral image of a difference image

Example: Calculate SSD for offset (-10, -15)

I.e. for a patch shifted -10 pixels down and -15 pixels to the right

• We can build an integral image of a difference image

Example: Calculate SSD for offset (-10, -15)

I.e. for a patch shifted -10 pixels down and -15 pixels to the right



Between this patch and this patch

• We can build an integral image of a difference image


Example: Calculate SSD for offset (-10, -15)



Between this patch and this patch

• We can build an integral image of a difference image

Example: Calculate SSD for offset (-10, -15)



Between this patch and this patch

• We can build an integral image of a difference image

Example: Calculate SSD for offset (-10, -15)



- 1. Calculate <u>per-pixel</u> squared difference for the entire image, offset with itself
- → Difference Image

• We can build an integral image of a difference image

Example: Calculate SSD for offset (-10, -15)



- 1. Calculate <u>per-pixel</u> squared difference for the entire image, offset with itself
- → Difference Image
- 2. Calculate Integral Image of the difference image
- → We can now read SSD *for all patches for this offset* in constant time!

• We can build an integral image of a difference image

Example: Calculate SSD for offset (-10, -15)



- 1. Calculate <u>per-pixel</u> squared difference for the entire image, offset with itself
- → Difference Image
- 2. Calculate Integral Image of the difference image
- → We can now read SSD *for all patches for this offset* in constant time!

Note: There are as many offset as there are pixels in the search window!

Problems

- Memory Overhead?
- Parallel Computation of the Integral Image?
- No benefit for smaller patch-sizes
- Some Distance Metrics not supported



Original



Noisy



Denoised





Closing Remarks

We have just implemented & understood the NL-Means filter:

$$NL(u(p)) = \frac{1}{C(p)} \int f(d(B(p), B(q))u(q)) dq$$

- Many other more complicated non-local methods:
 - BM3D/BM3D-SAPCA,
 - Non-Local Bayes,
 - PLOW, etc...



References

 ANTONI BUADES, BARTOMEU COLL, AND JEAN-MICHEL MOREL, Non-Local Means Denoising, Image Processing On Line, 1 (2011). http://dx.doi.org/10.5201/ipol.2011.bcm_nlm

• Gabriele Facciolo, Nicolas Limare, and Enric Meinhardt-Llopis, *Integral Images for Block Matching*, <u>Image Processing On Line</u>, 4 (2014), pp. 344–369. http://dx.doi.org/10.5201/ipol.2014.57